



OpenMP tasks



As we have seen in the previous example (`02_sections_nested_irregular.c`), it is sometimes possible to parallelize a workflow which is irregular or runtime-dependent using OpenMP sections.

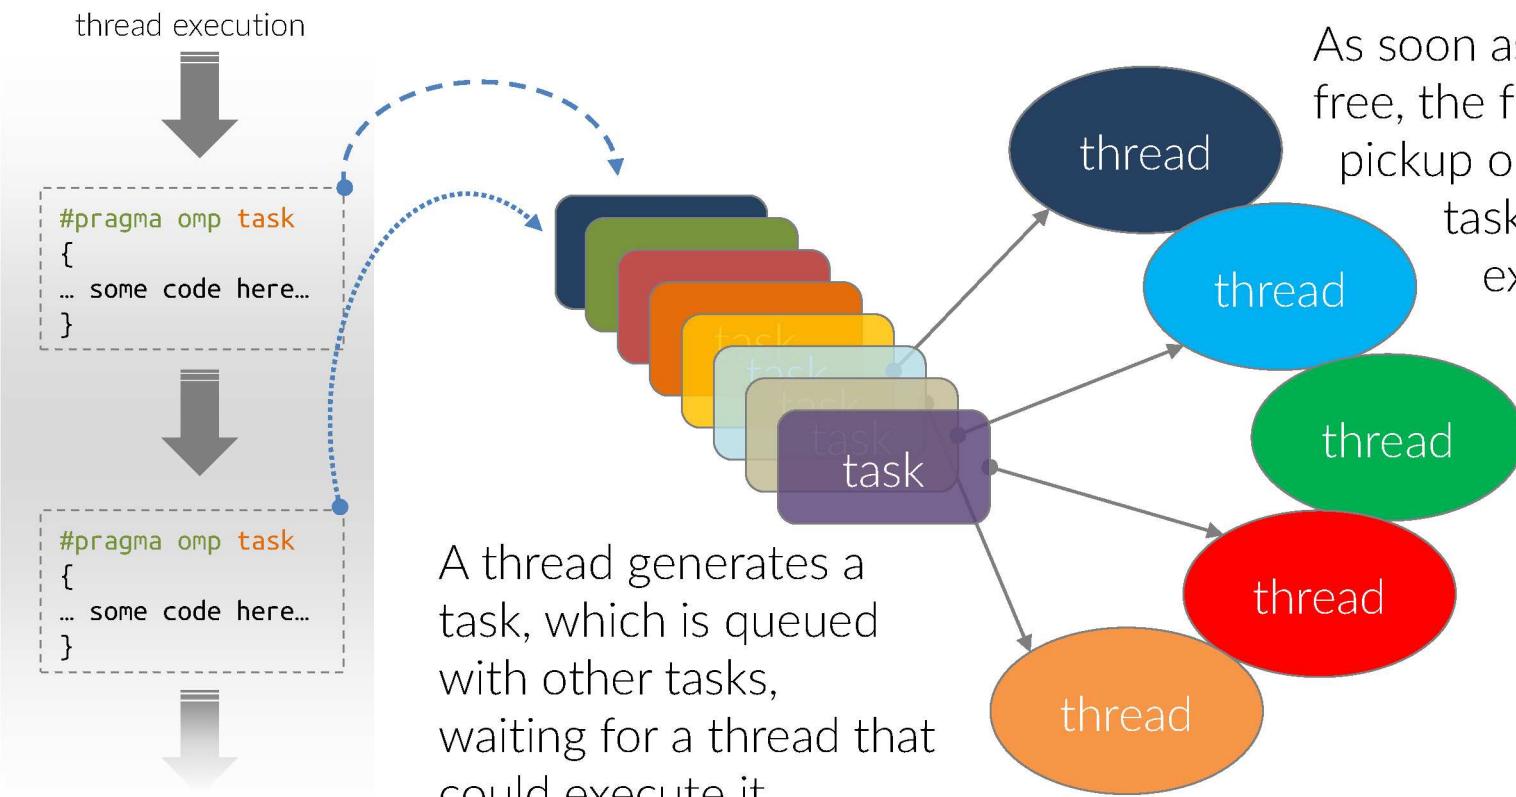
However, often the solution is quite ugly and convoluted.

Since version 3.0, OpenMP offers a new elegant construct designed for this class of problems: the OpenMP task.

What happens under the hood is that OpenMP creates a “bunch of work” along with the data and the local variables it needs, and schedules it for execution at some point in the future.



OpenMP tasks





OpenMP tasks



As almost everything else in OpenMP, a task must be generated *inside* a parallel region and it is linked to a specific block of code.

If its execution is not properly “protected”, it might be executed by *more* than one thread (i.e. by all threads that encounter the task definition), which is not in general what we want.

To guarantee that each task is executed only once, every task must be generated within a `single` or `master` region.

The `single` region is preferable because of its implied barrier that makes all tasks to be completed before passing. In case you use a `master` region, pay attention to the execution flow.

Moreover, the `master` has often the heavier burden so it’s best to user a `single` region, possibly with the `nowait` clause.



```
#pragma parallel region
{
    ...
    #pragma omp single nowait
    {
        while( !end_of_list(node) ) {
            if( node_is_to_be_processed(node) )
                #pragma omp task
                process_node ( node );
            node = next_node( node );
        }
        ...
    }
}
```

A classical example:
traversing a linked list

A task is generated for each
node that must be processed

The calling thread continues
traversing the linked list

Due to the `nowait` clause, all the threads skip
the implied barrier at the end of the `single`
region and wait here for being assigned a task



OpenMP tasks

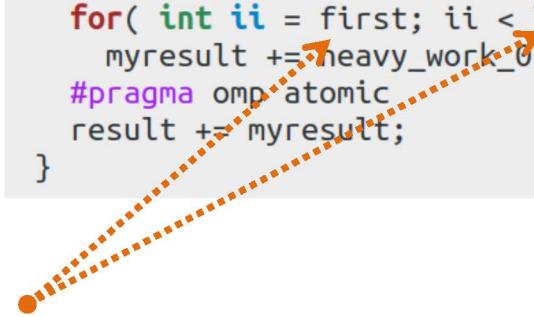


A second key point to account for when dealing with the asynchronous execution is the *data environment*.

A task is a confined code section that performs some operations on a data set, that is referred at the moment of the task creation.

You are in charge of ensuring that that reference will still be valid *at the moment of execution*, which is somewhere in the future.

```
#pragma omp task shared(result) untied
{
    double myresult = 0;
    for( int ii = first; ii < last; ii++)
        myresult += heavy_work_0(array[ii]);
    #pragma omp atomic
    result += myresult;
}
```



Both `first` and `last`, which are two shared variables, are key variables for the task execution.

What if they are keep changing?

At the moment of execution, their value could be different than at the moment of task creation, and then the processing would be totally different than the original intention.



OpenMP tasks



A second key point to account for when dealing with the asynchronous execution is the *data environment*.

A task is a confined code section that performs some operations on a data set, that is referred at the moment of the task creation.

You are in charge of ensuring that that reference will still be valid *at the moment of execution*, which is somewhere in the future.

The values of variables that are susceptible to change and that enter in the execution of the task must be protected to ensure the correctness of the task itself.

With the `firstprivate` clause, we are creating private local variables that will be referred to at the moment of the execution and will still have the correct value.

```
#pragma omp task firstprivate(first, last) shared(result) untied
{
    double myresult = 0;
    for( int ii = first; ii < last; ii++)
        myresult += heavy_work_0(array[ii]);
    #pragma omp atomic
    result += myresult;
}
```



OpenMP tasks



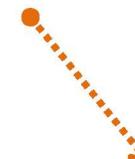
A second key point to account for when dealing with the asynchronous execution is the *data environment*.

A task is a confined code section that performs some operations on a data set, that is referred at the moment of the task creation.

You are in charge of ensuring that that reference will still be valid *at the moment of execution*, which is somewhere in the future.

With the `untied` clause, you are signalling that this task – if ever suspended – can be resumed by *any* free thread. The default is the opposite, a task to be tied to the thread that initially starts it.

If untied, you must take care of the data environment, of course: for instance, no `threadprivate` variables can be used, nor the thread number, and so on.



```
#pragma omp task firstprivate(first, last) shared(result) untied
{
    double myresult = 0;
    for( int ii = first; ii < last; ii++)
        myresult += heavy_work_0(array[ii]);
    #pragma omp atomic
    result += myresult;
}
```



OpenMP tasks synchronization



A third key point to catch with asynchronous execution, is about the *timing*, i.e. when a task is executed and how to synchronize them.

At the moment of creation, a task may be *deferred* or not, i.e. its execution may be scheduled for the future or immediately taken while the task region that has generated it is frozen.

There are some constructs that enforce synchronization:

barrier

Implicit or explicit barrier

taskwait

Wait on the completion of all child tasks of the current task

taskgroup

Wait on the completion of all child tasks of the current task **and** of their descendant



OpenMP tasks synchronization



```
#pragma omp parallel shared(result)
{
    double result1, result2, result3;

    #pragma omp single nowait
    {
        #pragma omp task shared(result1)
        {
            double myresult = 0;
            for( int jj = 0; jj < N; jj++ )
                myresult += heavy_work_0( array[jj] );
            #pragma omp atomic update
            result1 += myresult;
        }

        #pragma omp task shared(result2)
        {
            double myresult = 0;
            for( int jj = 0; jj < N; jj++ )
                myresult += heavy_work_1( array[jj] );
            #pragma omp atomic update
            result2 += myresult;
        }

        #pragma omp task shared(result3)
        {
            double myresult = 0;
            for( int jj = 0; jj < N; jj++ )
                myresult += heavy_work_2( array[jj] );
            #pragma omp atomic update
            result3 += myresult;
        }
    }

    #pragma omp taskwait

    #pragma omp atomic update
    result += result1;
    #pragma omp atomic update
    result += result2;
    #pragma omp atomic update
    result += result3;
}
```

You need the tasks to be completed *before* to arrive at this point where the final results are accumulated.

Due to the `nowait` clause, the implied barrier at the end of the `single` region is no respected and the threads are flowing freely beyond that region.

Without the `taskwait` directive, the threads that are not in the `single` region would execute the updates of `result` with meaningless data.



parallel_tasks/
03_tasks_wrong.c

finer implementation

parallel_tasks/
03_tasks_wrong.c





OpenMP tasks synchronization



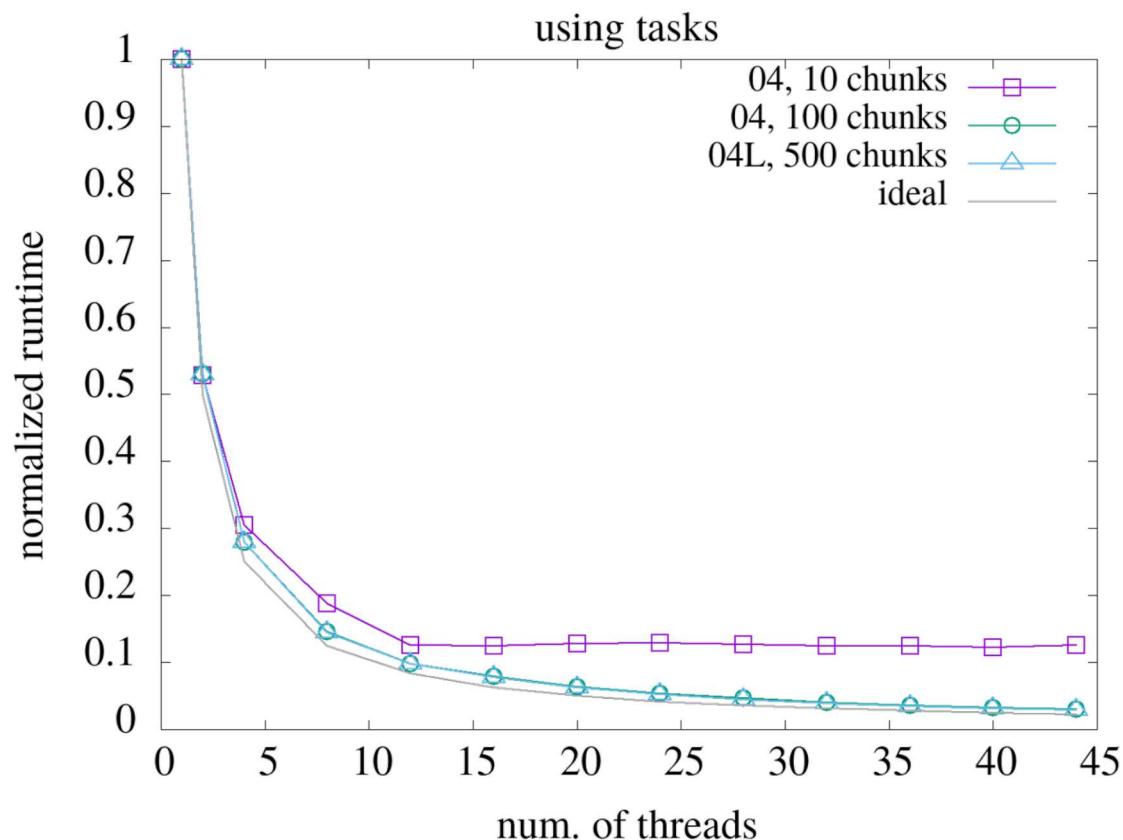
parallel_tasks/
03_tasks_wrong.c

- You need the tasks to be completed *before* to arrive at this point where the final results are accumulated.
- Due to the `nowait` clause, the implied barrier at the end of the `single` region is no respected and the threads are flowing freely beyond that region.
- Without the `taskwait` directive, the threads that are not in the `single` region would execute the updates of `result` with meaningless data.



Advanced
Parallelism

OpenMP...





OpenMP taskgroup



```
#pragma omp parallel proc_bind(close)
{
    #pragma omp single nowait
    {
        #pragma omp taskgroup task_reduction(+:result)
        {
            int idx = 0;
            int first = 0;
            int last = chunk;

            while( first < N )
            {
                last = (last >= N)?N:last;
                for( int kk = first; kk < last; kk++, idx++ )
                    array[idx] = min_value + lrand48() % max_value;

                #pragma omp task in_reduction(+:result) firstprivate(first, last) untied
                {
                    ...
                }
                #pragma omp task in_reduction(+:result) firstprivate(first, last) untied
                {
                    ...
                }
                #pragma omp task in_reduction(+:result) firstprivate(first, last) untied
                {
                    ...
                }
                first += chunk;
                last += chunk;
            }
        }
        #pragma omp taskwait
    } // close parallel region
}
```

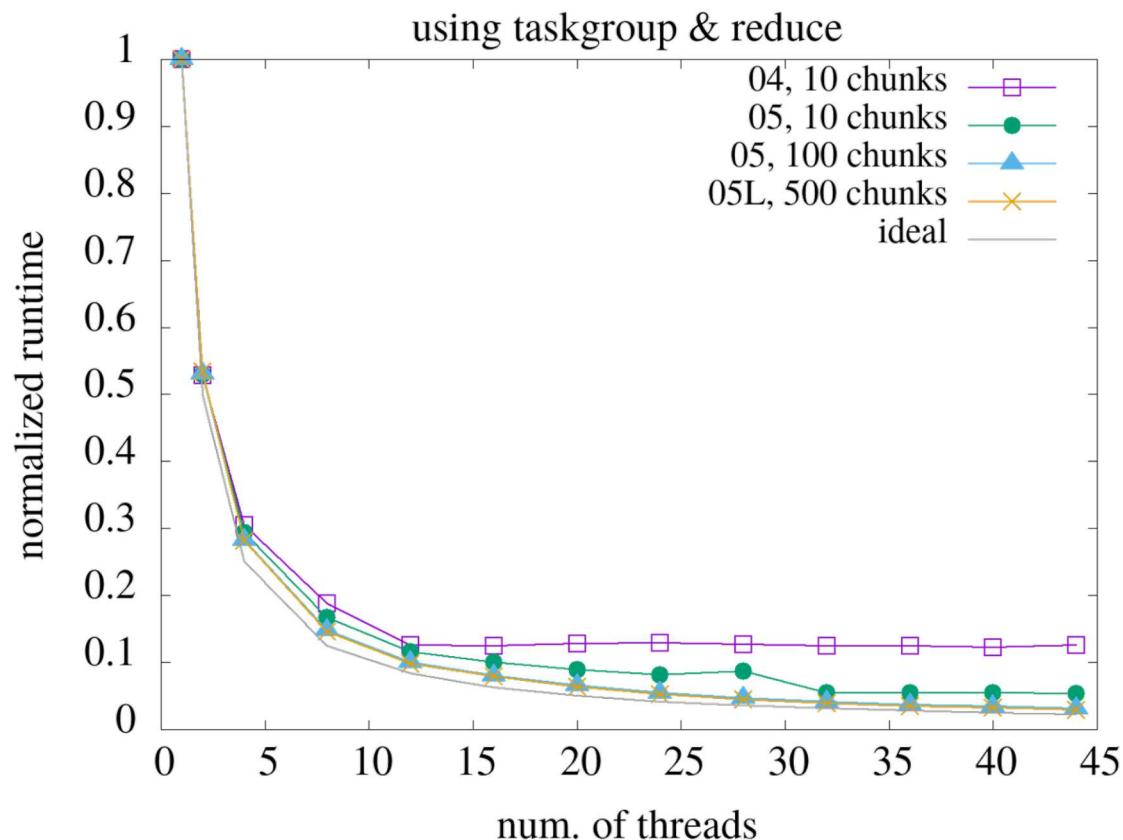
parallel_tasks/
05_task_taskgroup.c

A taskgroup region is declared: at its end, the completion of all tasks generated within it, and of their descendant, is explicitly ensured.

This task are participating to the reduction



OpenMP...





Advanced
Parallelism

OpenMP taskloop



```
#pragma omp parallel proc_bind(close)
{
    #pragma omp single nowait
    {
        //##pragma omp taskloop grainsize(N/1000) reduction(+:result)
        #pragma omp taskloop num_tasks(N/10) reduction(+:result)
        for( int ii = 0; ii < N; ii++ )
        {
            array[ii] = min_value + lrand48() % max_value;
            result += heavy_work_0(array[ii]) +
                heavy_work_1(array[ii]) +
                heavy_work_2(array[ii]);
        }
    }
    PRINTF("* initializer thread: initialization lasted %g seconds\n", CPU_TIME_th - tstart );
} // close parallel region

double tend = CPU_TIME;
#endif
```



parallel_tasks/
06_task_taskloop.c



OpenMP taskloop



```
#pragma omp parallel proc_bind(close)
{
    #pragma omp single nowait
    {
        //#pragma omp taskloop grainsize(N/1000) reduction(+:reslt)
        #pragma omp taskloop num_tasks(N/10) reduction(+:result)
        for( int ii = 0; ii < N; ii++ )
        {
            array[ii] = min_value + lrand48() % max_value;
            result += heavy_work_0(array[ii]) +
                heavy_work_1(array[ii]) +
                heavy_work_2(array[ii]);
        }
        PRINTF("* initializer thread: initialization lasted %g seconds\n", CPU_TIME_th - tstart );
    } // close parallel region

    double tend = CPU_TIME;
#endif
```

A taskloop region is declared:
it blends the flexibility of tasking
with the ease of loops

Tasks are created for each
iteration



parallel_tasks/
05_task_taskloop.c



OpenMP taskloop



```
#pragma omp parallel proc_bind(close)
{
    #pragma omp single nowait
    {
        //##pragma omp taskloop grainsize(N/1000) reduction(+:result)
        #pragma omp taskloop num_tasks(N/10) reduction(+:result)
        for( int ii = 0; ii < N; ii++ )
        {
            array[ii] = min_value + lrand48() % max_value;
            result += heavy_work_0(array[ii]) +
                heavy_work_1(array[ii]) +
                heavy_work_2(array[ii]);
        }
        PRINTF("* initializer thread: initialization lasted %g seconds\n", CPU_TIME_th - tstart );
    } // close parallel region

    double tend = CPU_TIME;
#endif
```

parallel_tasks/
05_task_taskloop.c

To limit overhead, you can control the task generation by using of num_tasks and grainsize clauses

Tasks are created for each iteration
Tasks are created accordingly to clauses